

# Lecture 001

## Statistical learning: Foundations

---

Edward Rubin  
05 January 2021

Admin

# Admin

## Today

### In-class

- *Course website:* <https://github.com/edrubin/EC524W21/>
- *Resources*
  - **RStudio** cheatsheets, books, and tutorials
  - **UO library**
  - See course page for more...
- Formalizing statistical learning, notation, goals (and problems)

↻ Eugene R Users Retweeted



**Ryann Crowley**

@ryann\_crowley



Interested in an intro to [#MachineLearning](#), sharing info, developing an ML community, or curious about ML opportunities in Eugene? Come meet/learn at "Machine learning for the web" hosted by the Eugene Web Developers [@EugeneRUsers](#) [@uodatasci](#) [@WiMLDS\\_PDX](#)

Tweet; h/t: Grant McDermott

# Admin

## Upcoming

### Readings

- *Today*
  - ISL Ch1–Ch2
  - **Prediction Policy Problems** by Kleinberg *et al.* (2015)
- *Next*
  - ISL Ch. 3–4

**Problem set** Soon.

# Statistical learning

# Statistical learning

## What is it?

**Statistical learning** is a **set of tools** developed **to understand/model data**.

## Examples

- **Regression analysis** quantifies the relationship between an outcome and a set of explanatory variables—most usefully in a causal setting.
- **Exploratory data analysis** (EDA) is a preliminary, often graphical, "exploration" of data to understand levels, variation, missingness, *etc.*
- **Classification trees** search through explanatory variables, splitting along the most "predictive" dimensions (random forests extend trees).
- **Regression trees** extend *classification trees* to numerical outcomes (random forests extend, as well).
- **K-means clustering** partitions observations into K groups (clusters) based upon a set of variables.

# Statistical learning

## What is it good for?

A lot of things. We tend to break statistical-learning into two(-ish) classes:

1. **Supervised learning** builds ("learns") a statistical model for predicting an **output** ( $\mathbf{y}$ ) given a set of **inputs** ( $\mathbf{x}_1, \dots, \mathbf{x}_p$ ), i.e., we want to build a model/function  $f$


$$\mathbf{y} = f(\mathbf{x}_1, \dots, \mathbf{x}_p)$$

that accurately describes  $\mathbf{y}$  given some values of  $\mathbf{x}_1, \dots, \mathbf{x}_p$ .

2. **Unsupervised learning** learns relationships and structure using only **inputs** ( $\mathbf{x}_1, \dots, \mathbf{x}_p$ ) without any *supervising* output—letting the data "speak for itself."



**Semi-supervised learning** falls somewhere between these supervised and unsupervised learning—generally applied to supervised tasks when labeled **outputs** are incomplete.



**Supervised Learning**

**Unsupervised Learning**

Source

# Statistical learning

## Output

We tend to further break **supervised learning** into two groups, based upon the **output** (the **outcome** we want to predict):

1. **Classification tasks** for which the values of **y** are discrete categories  
*E.g.*, race, sex, loan default, hazard, disease, flight status
2. **Regression tasks** in which **y** takes on continuous, numeric values.  
*E.g.*, price, arrival time, number of emails, temperature

*Note*<sub>1</sub> The use of *regression* differs from our use of *linear regression*.

*Note*<sub>2</sub> Don't get tricked: Not all numbers represent continuous, numerical values—*e.g.*, zip codes, industry codes, social security numbers.<sup>†</sup>

† Q Where would you put responses to 5-item Likert scales?

# Statistical learning

## *Why Learning?*

**Q** What puts the "learning" in statistical/machine learning?

**A** Most learning models/algorithms will **tune model parameters** based upon the observed dataset—learning from the data.

# Notation

Our class will typically follow the notation and definitions of *ISL*.

# Notation

## Data

$n$  gives the number of observations

$p$  represents the number of variables available for predictions

$\mathbf{X}$  is our  $n \times p$  matrix of predictors

- Other names **features**, inputs, independent/explanatory variables, ...
- $x_{i,j}$  is observation  $i$  (in  $1, \dots, n$ ) on variable  $j$  (for  $j$  in  $1, \dots, p$ )

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & x_{1,2} & \cdots & x_{1,p} \\ x_{2,1} & x_{2,2} & \cdots & x_{2,p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n,1} & x_{n,2} & \cdots & x_{n,p} \end{bmatrix}$$

# Notation

## Dimensions of $\mathbf{X}$

Now let us split our  $\mathbf{X}$  matrix of predictors by its two dimensions.

**Observation**  $i$  is a  $p$ -length vector

$$\mathbf{x}_i = \begin{bmatrix} x_{i,1} \\ x_{i,2} \\ \vdots \\ x_{i,p} \end{bmatrix}$$

**Variable**  $j$  is a  $n$ -length vector

$$\mathbf{x}_j = \begin{bmatrix} x_{1,j} \\ x_{2,j} \\ \vdots \\ x_{n,j} \end{bmatrix}$$

Applied to R:

- `dim(x_df) =  $n$   $p$`
- `nrow(x_df) =  $n$ ; ncol(x_df) =  $p$`
- `x_df[1,]` ( $i = 1$ ); `x_df[,1]` ( $j = 1$ )

# Notation

## Outcomes

In supervised settings, we will denote our **outcome variable** as  $\mathbf{y}$ .

*Synonyms* output, outcome, dependent/response variable, ...

The **outcome** for our  $i^{\text{th}}$  observation is  $y_i$ . Together the  $n$  observations form

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

and our full dataset is composed of  $\left\{ (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n) \right\}$ .



Back to the problem of (supervised) statistical learning...

# Statistical learning

## The goal

As defined before, we want to *learn* a model to understand our data.

1. Take our (numeric) **output**  $\mathbf{y}$ .
2. Imagine there is a **function**  $f$  that takes **inputs**  $\mathbf{X} = \mathbf{x}_1, \dots, \mathbf{x}_p$  and maps them, plus a random, mean-zero **error term**  $\varepsilon$ , to the **output**.

$$\mathbf{y} = f(\mathbf{X}) + \varepsilon$$

**Q** What is  $f$ ?

**A** *ISL*:  $f$  represents the *systematic* information that  $\mathbf{X}$  provides about  $\mathbf{y}$ .

**Q** How else can you describe  $f$ ?

# Statistical learning

## Our missing $f$

$$\mathbf{y} = f(\mathbf{X}) + \varepsilon$$

Q  $f$  is unknown (as is  $\varepsilon$ ). What should we do?

A Use the observed data to learn/estimate  $f(\cdot)$ , i.e., construct  $\hat{f}$ .<sup>†</sup>

Q Okay. How?

A *How do I estimate  $f$ ?* is one way to phrase *all questions* that underly statistical learning—model selection, cross validation, evaluation, etc.

All of the techniques, algorithms, tools of stat. learning attempt to accurately recover  $f$  based upon the settings' goals/limitations.

You'll have to wait on any real/specific answers...

<sup>†</sup> More notation: hats (^) are estimators/estimates.

# Statistical learning

## Learning from $\hat{f}$

There are two main reasons we want to learn about  $f$

1. **Causal inference settings** How do changes in  $\mathbf{X}$  affect  $\mathbf{y}$ ?

The territory of EC523 and EC525.

2. **Prediction problems** Predict  $\mathbf{y}$  using our estimated  $f$ , i.e.,

$$\hat{\mathbf{y}} = \hat{f}(\mathbf{X})$$

our *black-box setting* where we care less about  $f$  than  $\hat{\mathbf{y}}$ .<sup>†</sup>

Similarly, in causal-inference settings, we don't particularly care about  $\hat{\mathbf{y}}$ .

<sup>†</sup> You shouldn't actually treat your prediction methods as total black boxes.

# Statistical learning

## Prediction errors

As tends to be the case in life, you will make errors in predicting  $\mathbf{y}$ .

The accuracy of  $\hat{\mathbf{y}}$  depends upon **two errors**:

1. **Reducible error** The error due to  $\hat{f}$  imperfectly estimating  $f$ .  
*Reducible* in the sense that we could improve  $\hat{f}$ .
2. **Irreducible error** The error component that is outside of the model  $f$ .  
*Irreducible* because we defined an error term  $\epsilon$  unexplained by  $f$ .

*Note* As its name implies, you can't get rid of *irreducible* error—but we can try to get rid of *reducible* errors.

# Statistical learning

## Prediction errors

Why we're stuck with *irreducible* error

$$\begin{aligned} E\left[\{\mathbf{y} - \hat{\mathbf{y}}\}^2\right] &= E\left[\left\{f(\mathbf{X}) + \varepsilon + \hat{f}(\mathbf{X})\right\}^2\right] \\ &= \underbrace{\left[f(\mathbf{X}) - \hat{f}(\mathbf{X})\right]^2}_{\text{Reducible}} + \underbrace{\text{Var}(\varepsilon)}_{\text{Irreducible}} \end{aligned}$$

In less math:

- If  $\varepsilon$  exists, then  $\mathbf{X}$  cannot perfectly explain  $\mathbf{y}$ .
- So even if  $\hat{f} = f$ , we still have irreducible error.

Thus, to form our **best predictors**, we will **minimize reducible error**.

# Statistical learning

## Which type of $\hat{f}$ ?

Once you have your **inputs** ( $\mathbf{X}$ ) and **output** ( $\mathbf{y}$ ) data, you still need to decide how parametric your  $\hat{f}$  should be.<sup>†</sup>

**Parametric methods** assume a function typically involve two steps

1. Select a functional form (shape) to represent  $f$
2. Train your selected model on your data  $\mathbf{y}$  and  $\mathbf{X}$ .

**Non-parametric methods** avoid explicit assumption about the shape of  $f$ . Attempt to **flexibly fit** the data, while trying to **avoid overfitting**.

<sup>†</sup> I'm saying "how parametric" b/c some methods are much more parametric than others.

# Statistical learning

## Which type of $\hat{f}$ ?

Methods' parametric assumptions come with tradeoffs.

### **Parametric methods**

- + Simpler to estimate and interpret.
- If assumed functional form is bad, model performance will suffer.

### **Non-parametric methods**

- + Fewer assumptions. More flexibility.
- Lower interpretability. Susceptible to overfitting. Want lots of data.



**Example:** Let's start with a pretty funky, nonlinear function.

**Truth:** The (nonlinear)  $f(\mathbf{X})$  that we hope to recover.

**The sample:**  $n = 70$  randomly drawn observations for  $\mathbf{y} = f(\mathbf{x}_1, \mathbf{x}_2) + \varepsilon$

**Estimated linear-regression model:**  $\hat{\mathbf{y}} = \hat{\beta}_0 + \hat{\beta}_1 \mathbf{x}_1 + \hat{\beta}_2 \mathbf{x}_2 + \hat{\beta}_3 \mathbf{x}_1 \mathbf{x}_2$

**Prediction error** from our fitted linear regression model

**k-nearest neighbors** (kNN) using  $k=5$  (a *non-parametric* method)

**k-nearest neighbors** (kNN) using  $k=10$  (notice increased smoothness)

**k-nearest neighbors** (kNN) using  $k=1$  (notice decreased smoothness)



**Prediction error** from our fitted kNN ( $k=5$ ) model

**Prediction error** from our fitted kNN ( $k=10$ ) model

**Prediction error** from our fitted kNN ( $k=1$ ) model

Recall **Prediction error** from our fitted linear regression model

# Model accuracy

## Questions

1. Which of the methods was the most flexible? Inflexible?
2. Why do you think kNN with  $k=1$  had such low prediction error?
3. How could we (better) assess model/predictive performance?
4. Why would we ever want to choose a less flexible model?

# Model accuracy

## Measurement

You probably will not be surprised to know that there is no one-size-fits-all solution in statistical learning.

**Q** How do we choose between competing models?

**A** We're a few steps away, but before we do anything, we need a way to **define model performance**.

# Model accuracy

## Subtlety

Defining performance can actually be quite tricky...

*Regression setting, 1* Which do you prefer?

1. Lots of little errors and a few really large errors.
2. Medium-sized errors for everyone.

*Regression setting, 2* Is a 1-unit error (e.g., \$1,000) equally bad for everyone?

# Model accuracy

## Subtlety

Defining performance can actually be quite tricky...

*Classification setting, 1* Which is worse?

1. False positive (*e.g.*, incorrectly diagnosing cancer)
2. False negative (*e.g.*, missing cancer)

*Classification setting, 2* Which is more important?

1. True positive (*e.g.*, correct diagnosis of cancer)
2. True negative (*e.g.*, correct diagnosis of "no cancer")



# Model accuracy

## MSE

**Mean squared error (MSE)** is the most common<sup>†</sup> way to measure model performance in a regression setting.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n \left[ y_i - \hat{f}(x_i) \right]^2$$

Recall:  $y_i - \hat{f}(x_i) = y_i - \hat{y}_i$  is our prediction error.

Two notes about MSE

1. MSE will be (relatively) very small when **prediction error** is nearly zero.
2. MSE **penalizes** big errors more than little errors (the squared part).

<sup>†</sup> *Most common* does not mean best—it just means lots of people use it.

# Model accuracy

## Training or testing?

Low MSE (accurate performance) on the data that trained the model isn't actually impressive—maybe the model is just overfitting our data.<sup>†</sup>

*What we want:* How well does the model perform **on data it has never seen**?

This introduces an important distinction:

1. **Training data:** The observations  $(y_i, x_i)$  used to **train** our model  $\hat{f}$ .
2. **Testing data:** The observations  $(y_0, x_0)$  that our model has yet to see—and which we can use to evaluate the performance of  $\hat{f}$ .

**Real goal: Low test-sample MSE** (not the training MSE from before).

<sup>†</sup> Recall the kNN performance for  $k=1$ .

**Next time:** model performance, the variance-bias tradeoff, and kNN

# Sources

These notes draw upon

- [An Introduction to Statistical Learning \(ISL\)](#)  
James, Witten, Hastie, and Tibshirani
- [Python Data Science Handbook](#)  
Jake VanderPlas

I pulled the comic from [Twitter](#).

# Table of contents

## Admin

- Today
- Upcoming

## Statistical learning

- Definition
- Classes

## Notation

- Source
- Data
- Dimensions of  $\mathbf{X}$
- Outcomes

## Statistical learning, continued

- The goal
- Prediction
- Parameterization

## Example

- Data-generating process (truth)
- Regression model
- kNN model

## Model accuracy

- Questions
- Subtlety
- MSE
- Training vs. testing

## Other

- Sources/references